

# IMAGE CLASSIFICATION WITH DEEP LEARNING

From Simple CNNs to Advanced Transfer Learning

Julia Parnis

CIFAR-10 IMAGE CLASSIFICATION  
PROJECT

**Attribution:** While conducted in a team setting, **all experiments and evaluations presented here were performed by me.**

# Project Overview

- ✓ **Goal:** Build an accurate image classification model for CIFAR-10 dataset
- ✓ **Dataset:** 60,000 images across 10 classes (32×32 RGB)
- ✓ **Challenge:** Improve accuracy through architecture design and optimization
- ✓ **Approach:** Systematic progression from simple to complex models



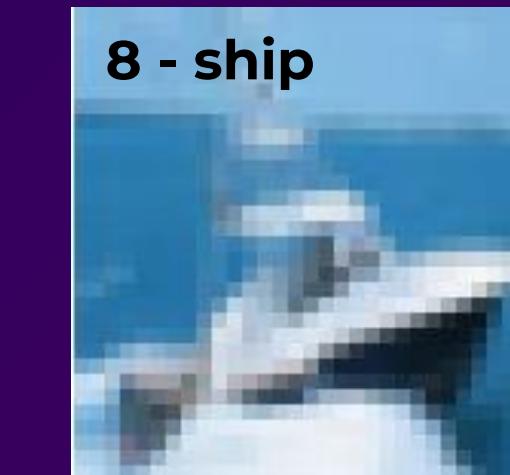
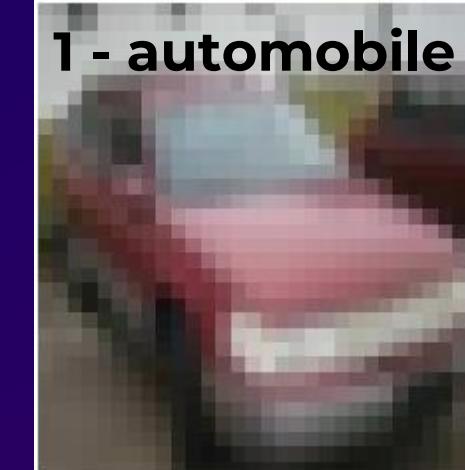
# CIFAR-10 dataset

- **Total:** 60,000 images
- **Training dataset:** 50,000 images
- **Test dataset:** 10,000 images
- **Image dimensions:**
  - 32x32 pixels,
  - 3 channels (RGB)
- **10 classes:** 0-9

- Even distribution of images per each class:
  - 5,000 images per class

## PREPROCESSING:

- Data was normalized:
  - $X_{train} = X_{train}/255$
  - $X_{test} = X_{test}/255$



# Our Approach

- 1 **Baseline CNN**  
Basic CNN model
- 2 **VGG Style Architecture (10 layers)**  
Deeper Feature Extraction
- 3 **ResNet-20 (20 layers)**  
Skip connections
- 4 **Transfer Learning**
  - EfficientNetV2





1

# Baseline CNN



Basic CNN model

Input (32x32x3)

# Baseline CNN Model

Input



Conv  $3 \times 3$ , 32 + ReLU

MaxPool  $2 \times 2$



Conv  $3 \times 3$ , 64 + ReLU

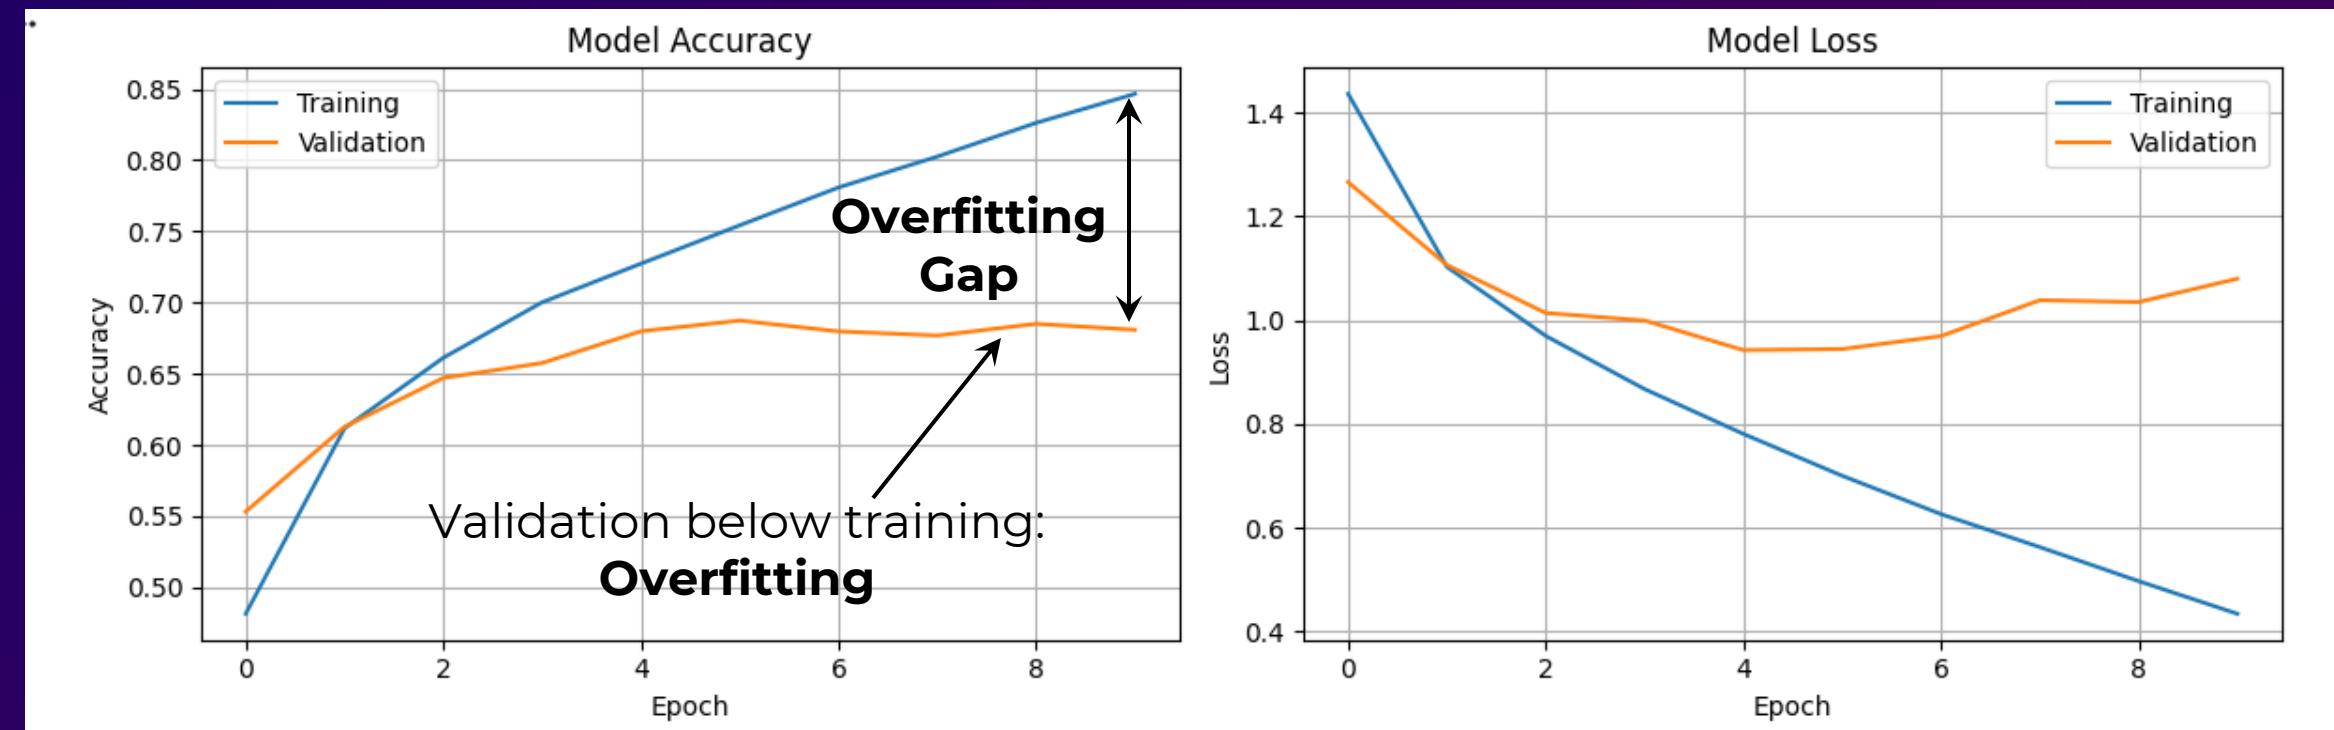
MaxPool  $2 \times 2$



Flatten/ Dense, 128

Dense 10 + Softmax

Model	Total Parameters	Key Modifications	Test Acc	Test Loss	Overfitting Gap	Epochs
5-layer CNN	316K	Dropout 0.2	67.5%	1.1	+17.8%	10



- The training and the validation datasets develop at different rates → **Severe Overfitting**
- Large test loss and low accuracy → **Further model improvement needed**



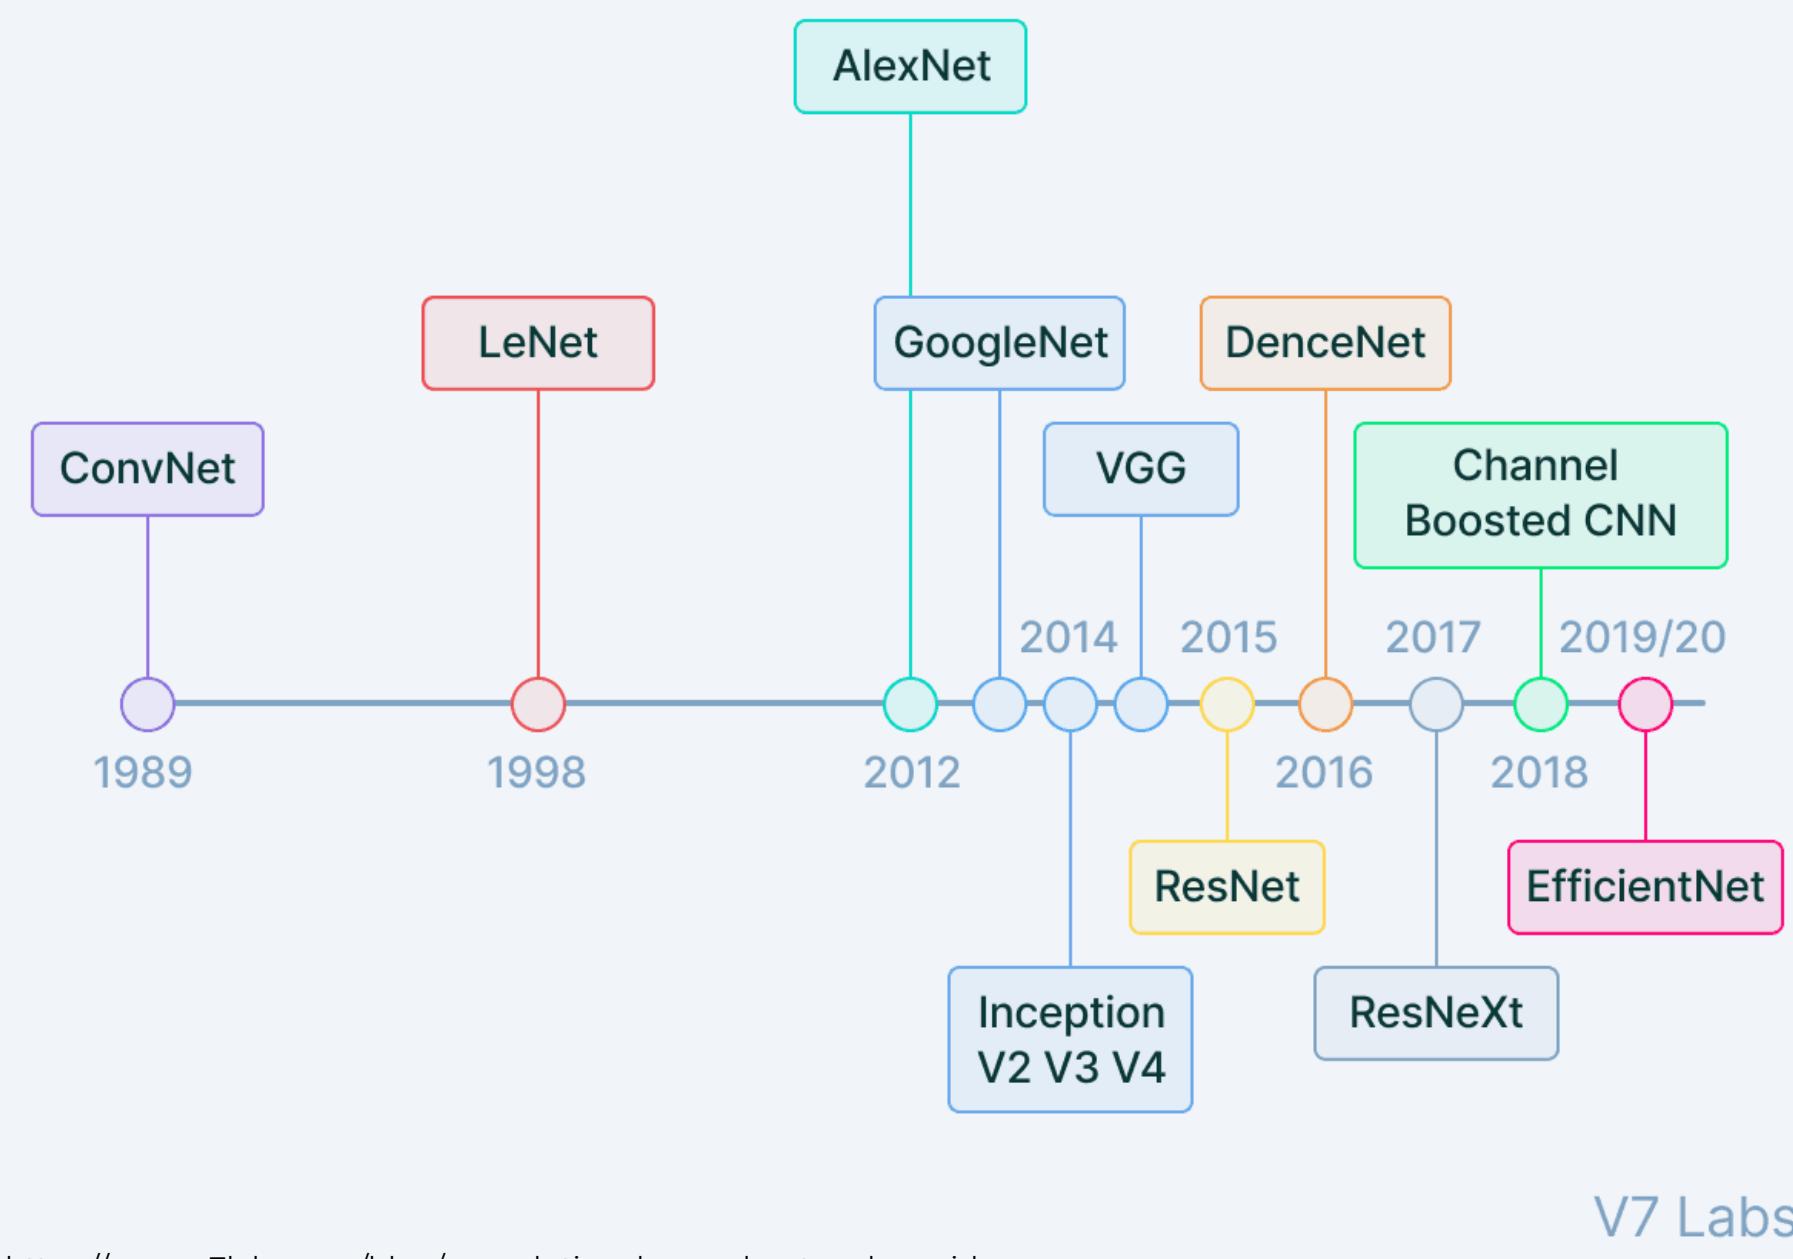
2

# VGG Style Architecture (10 layers)

Deeper Feature Extraction

# Why VGG Model is so Popular?

## Timeline of CNN models' Development



**VGG - (Visual Geometry Group, Oxford, 2014)**

**Key idea:** Go deeper using a simple, repeatable design

- **Small filters:** uses  $3 \times 3$  convolution filters
- **Block design:** Conv → Conv → MaxPool, repeat
- **Simple & consistent:** same building blocks throughout → easy to implement, and modify.
- **Practical impact:** became a popular **baseline and feature extractor** for many vision tasks

## Our tweaks:

- ✓ Callbacks (Early stop, ReduceLROnPlateau)
- ✓ Data augmentation
- ✓ Dropout

Input (32x32x3)

Input

Conv 3 x 3, 32 + ReLU (same)

Conv 3 x 3, 32 + ReLU (same)

MaxPool 2 x 2

Conv 3 x 3, 64 + ReLU (same)

Conv 3 x 3, 64 + ReLU (same)

MaxPool 2 x 2

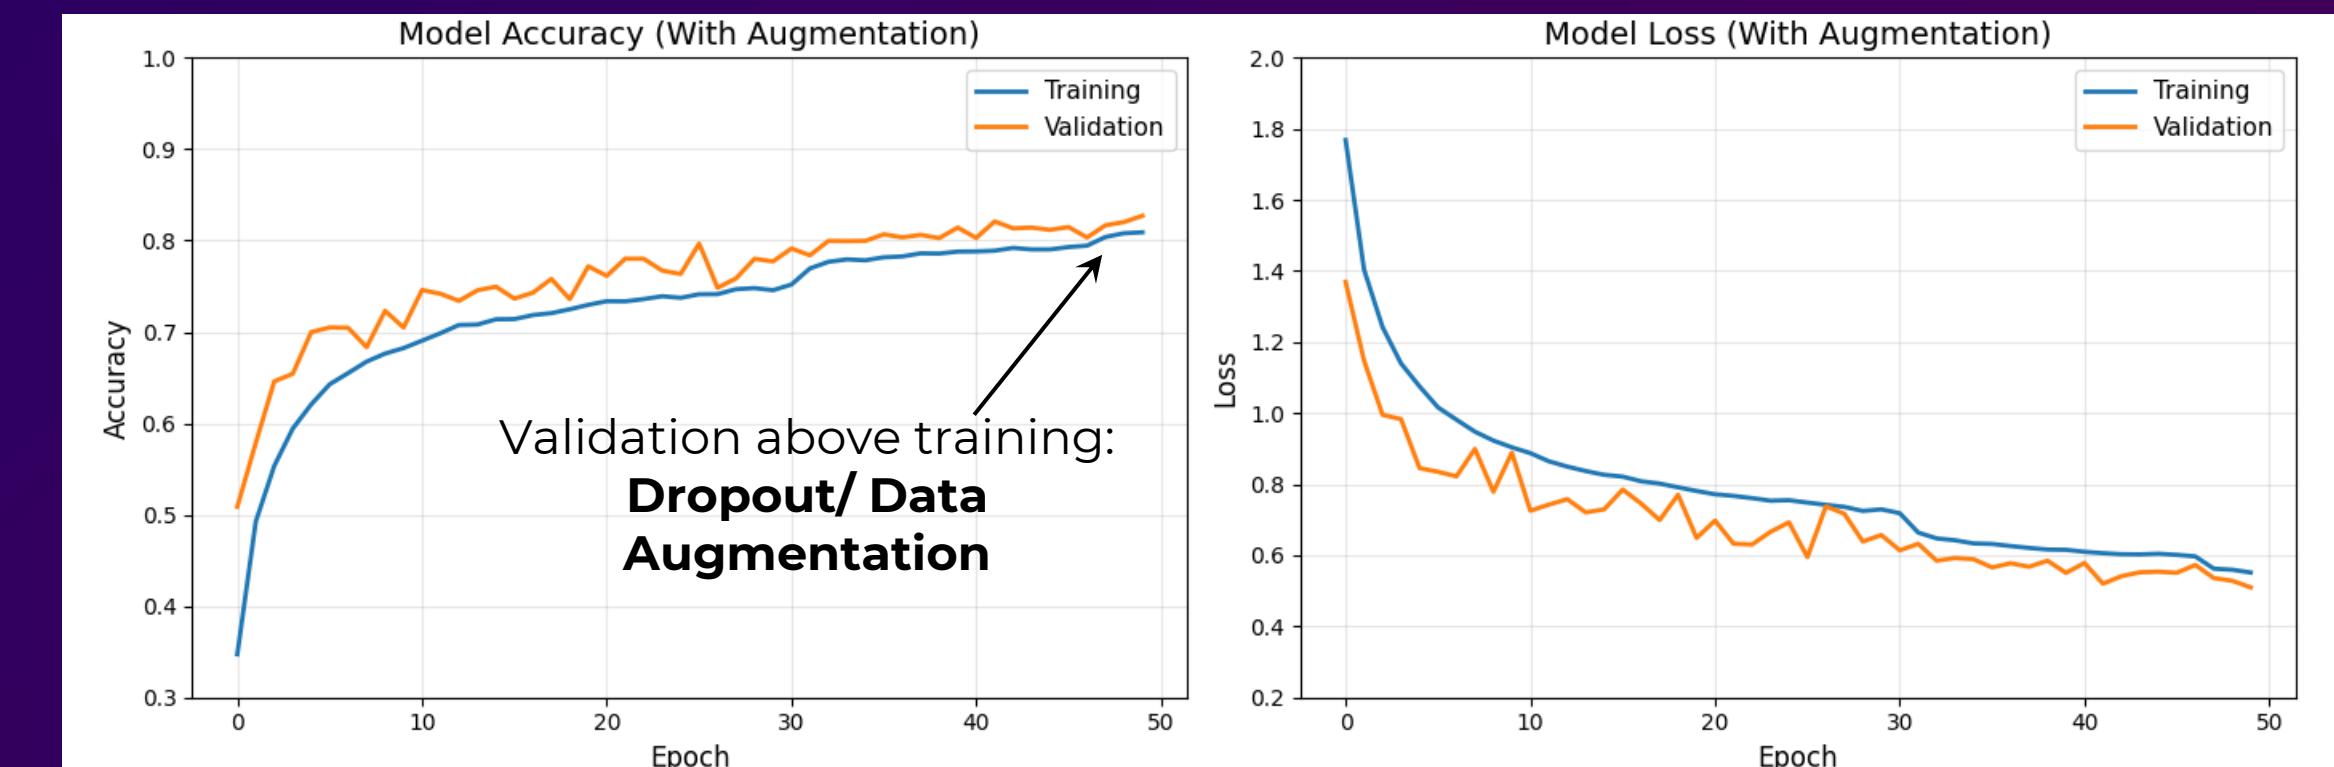
Flatten/ Dense, 128

Dense 10 + Softmax

- The training and the validation datasets develop together → **Model is learning efficiently**
- Deeper architecture + change in learning rate + data augmentation gradually improved model efficiency

# VGG-10 CNN: Deeper with Optimizations

Model	Total Parameters	Key Modifications	Test Acc	Test Loss	Overfitting Gap	Epochs
5-layer CNN	316K	Dropout 0.2	67.5%	1.1	+17.8%	10
VGG10-base	403K	Dropout 0.2	78.6%	0.6	+2.8%	20
VGG10+callbacks	403K	+ early stop, ReduceLR	80.5%	0.6	+7.4%	25/50
VGG10-optimized	403K	All+Data Augmentation	<b>82.1%</b>	0.5	-1.8%	50/50





3

# ResNet-20 (20 layers)



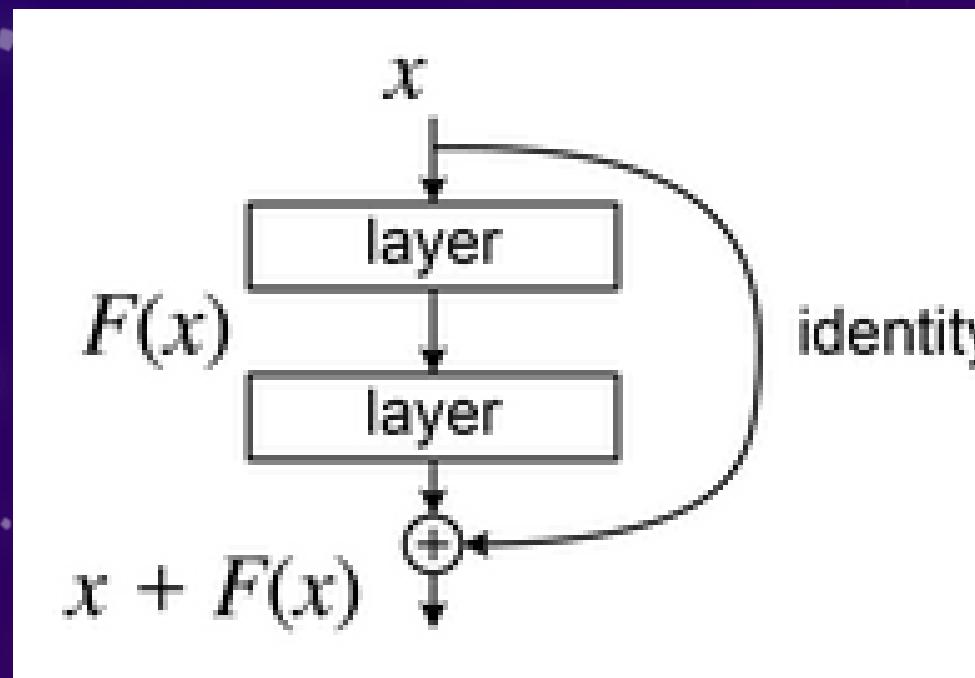
Skip connections

# ResNet-20: Solving the Depth Problem

## ResNet (2015) - “Residual Learning”

**Key idea:** skip connections learn a residual  $F(x)$  so the block outputs  $x + F(x)$

- **Paper:** Deep Residual Learning for Image Recognition (He et al., 2015)
- **Problem:** very deep nets started to train worse (“degradation” / vanishing gradients)



```
# Save input for skip connection
shortcut = x

# First conv layer
x = Conv2D(filters, (3,3), strides=stride, padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)

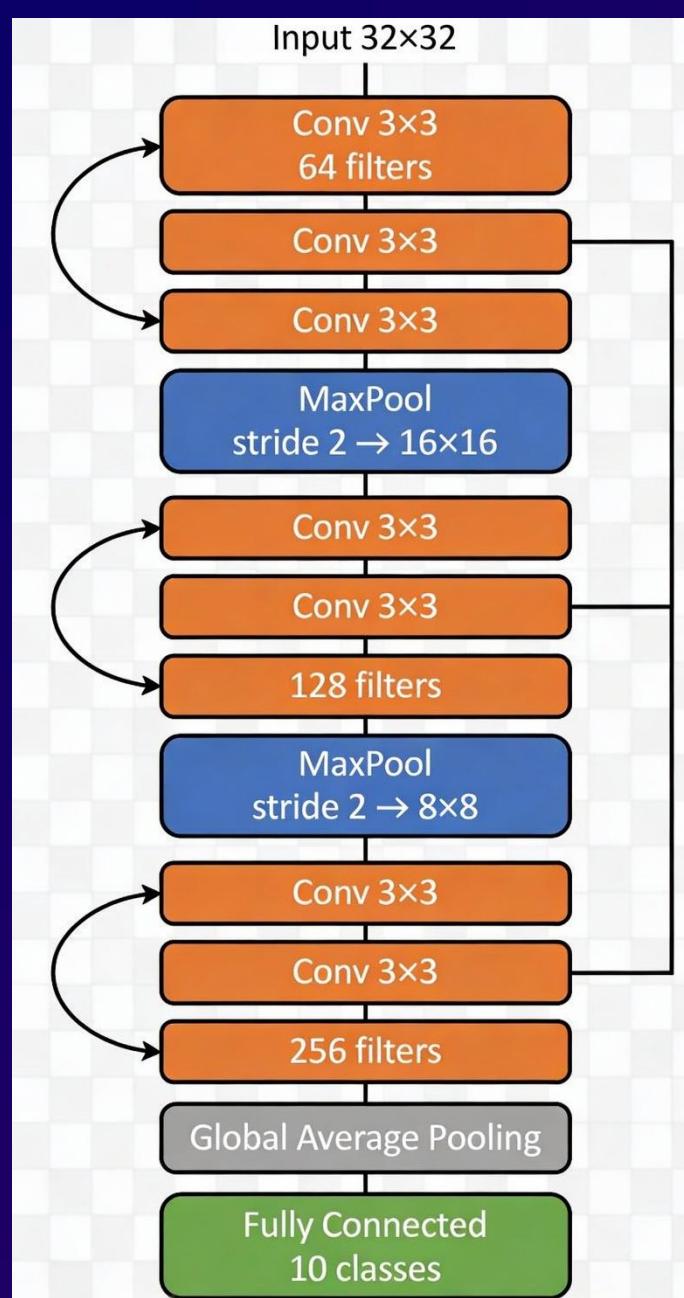
# Second conv layer
x = Conv2D(filters, (3,3), strides=1, padding='same')(x)
x = BatchNormalization()(x)

# Adjust shortcut if dimensions changed
if stride != 1 or shortcut.shape[-1] != filters:
    shortcut = Conv2D(filters, (1,1), strides=stride, padding='same')(shortcut)
    shortcut = BatchNormalization()(shortcut)

# Add skip connection
x = Add()([x, shortcut])
x = Activation('relu')(x)

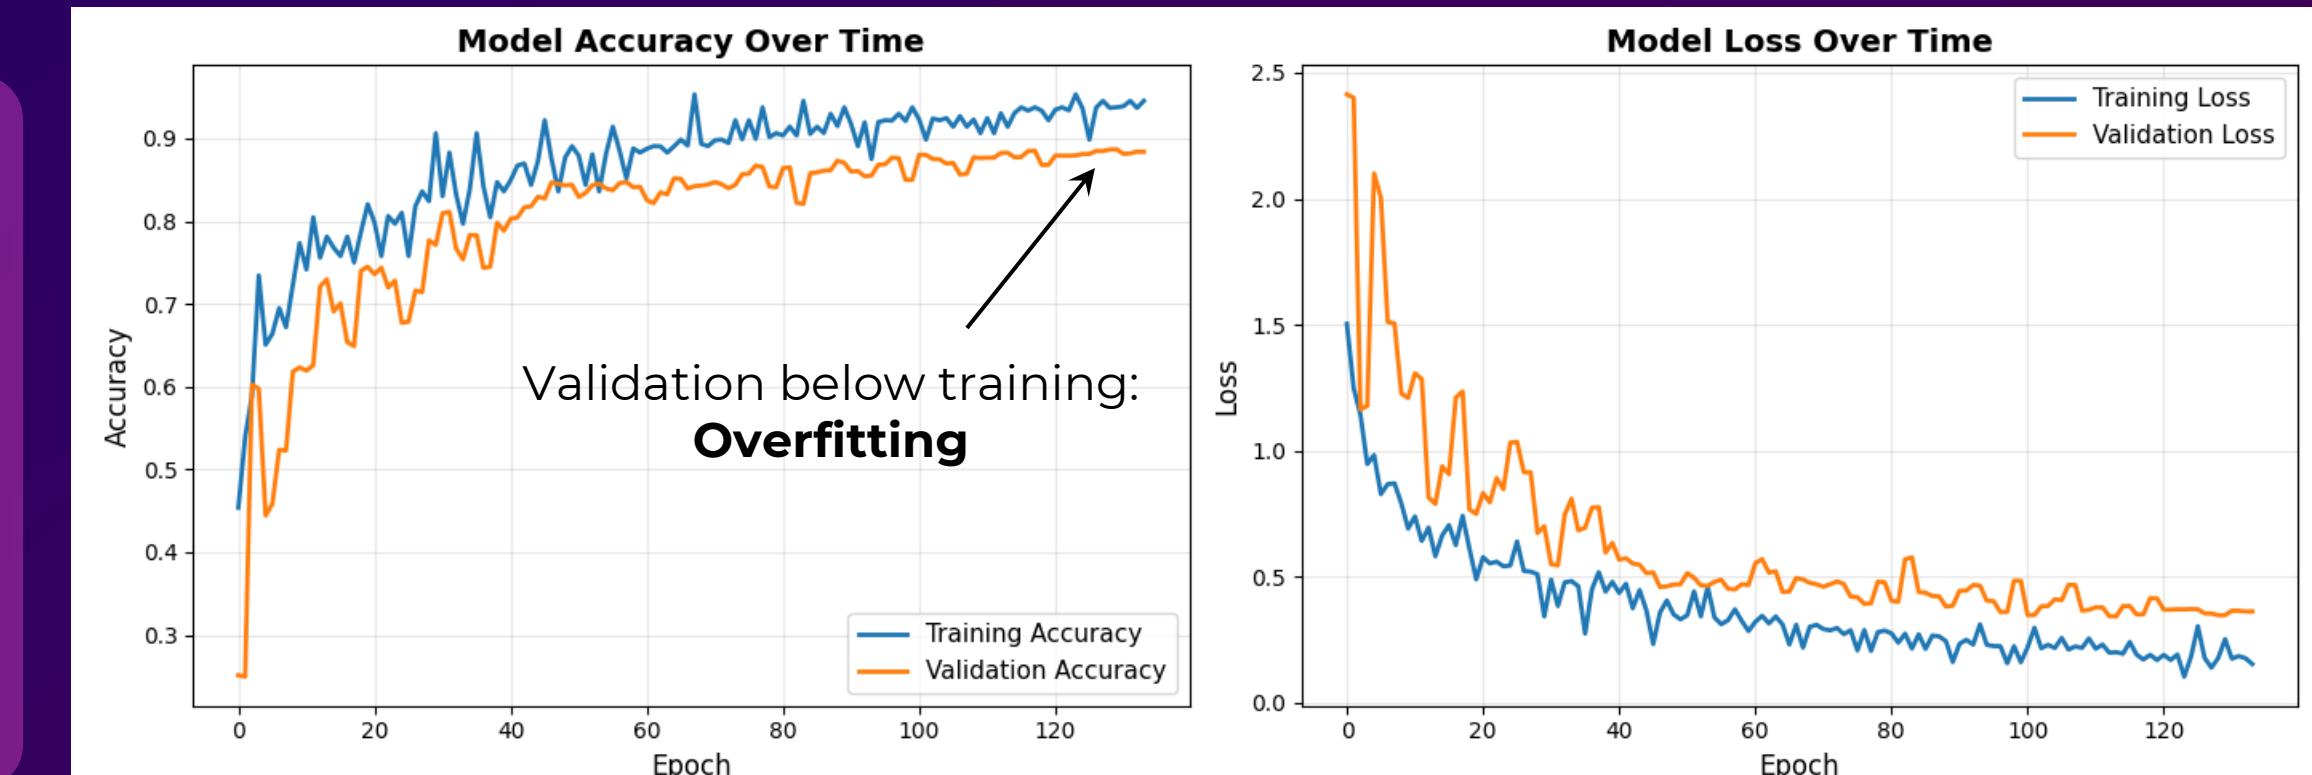
return x
```

# ResNet-20 CNN Model

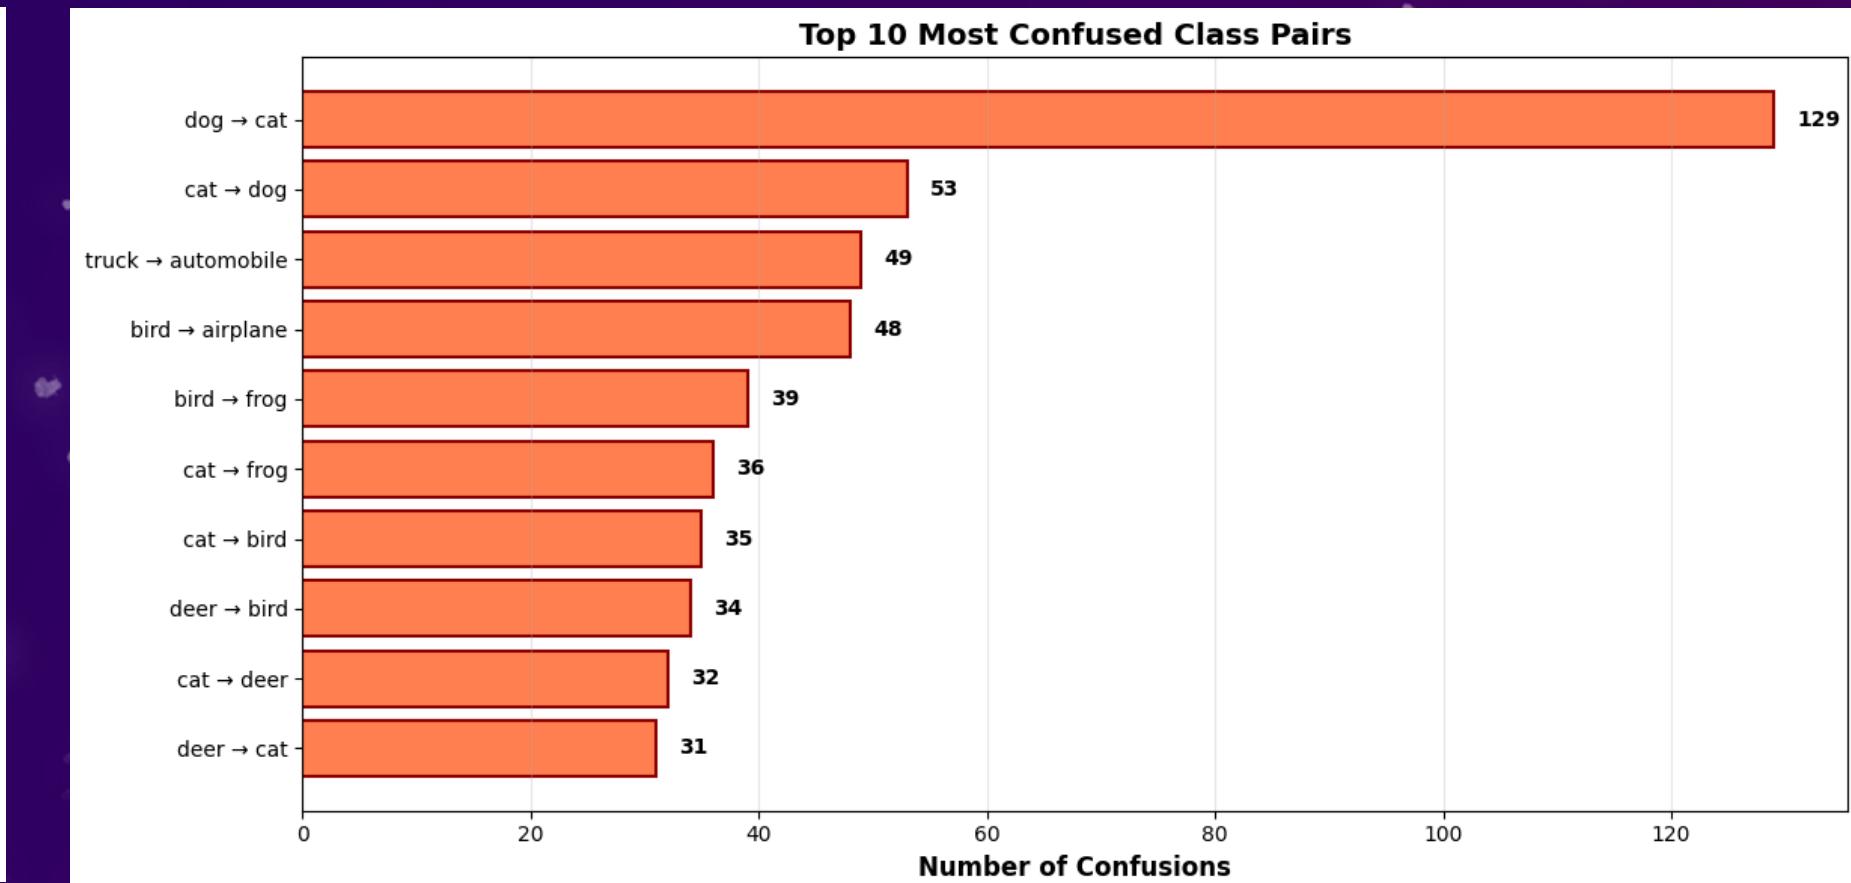


Model	Total Parameters	Key Modifications	Test Acc	Test Loss	Overfitting Gap	Epochs
5-layer CNN	316K	Dropout 0.2	67.5%	1.1	+17.8%	10
VGG10-base	403K	Dropout 0.2	78.6%	0.6	+2.8%	20
VGG10+callbacks	403K	+ early stop, ReduceLR	80.5%	0.6	+7.4%	25/50
VGG10-optimized	403K	All+Data Augmentation	82.1%	0.5	-1.8%	50/50
ResNet-20	275K	All previous, no Dropout	<b>87.7%</b>	0.4	+3.1%	114/200

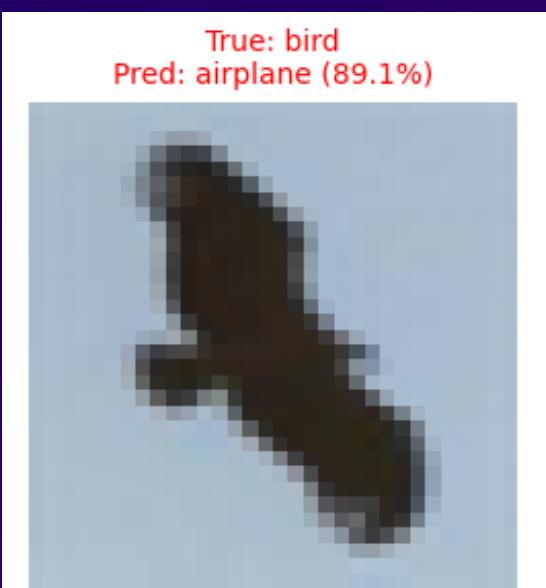
- The training and the validation datasets develop together → **Model is learning efficiently**
- Skip connections approach: ResNet-20 further improved model efficiency



# ResNet-20: Evaluation Dashboard



## Misclassified Examples: Where the model went wrong?





4

# Transfer Learning



EfficientNetV2

# EfficientNet: Smarter scaling for better accuracy–compute trade-off

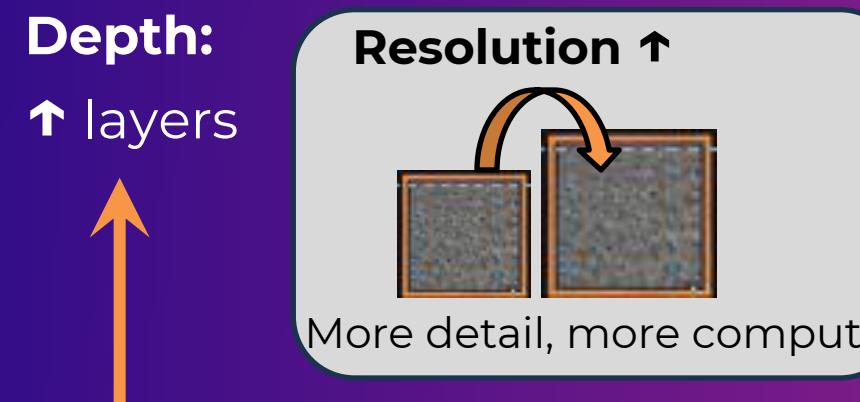
## EfficientNet (2019) – Core idea

Tan & Le (2019). "EfficientNet: Rethinking Model Scaling for CNNs." ICML 2019

- **Goal:** high accuracy with less compute (fewer parameters / fewer FLOPs)
- **Key idea:** compound scaling (scale depth + width + resolution together)
- **B0–B7:** B0 is the baseline; B1–B7 are scaled versions using the same rule

## Compound Scaling:

Depth + Width + Resolution



Stronger accuracy – efficiency balance

## EfficientNetV2 (2021): faster training



- **Fused blocks:** faster early layers



- **Progressive learning:** small → bigger images



- **Training speed:** faster training overall

## Transfer Learning Pipeline

**Input:** 32 x 32 x 3

**Resize:** 224 x 224 x 3

### Phase 1: Feature extraction

Freeze backbone – train only head

### Backbone:

EfficientNetV2B0 (ImageNet)

### Pooling:

Global Average Pool

### Phase 2 Fine-tuning

Unfreeze last N layers – train

### Head:

BN → Dropout → Dense(256) → BN → Dropout

### Output:

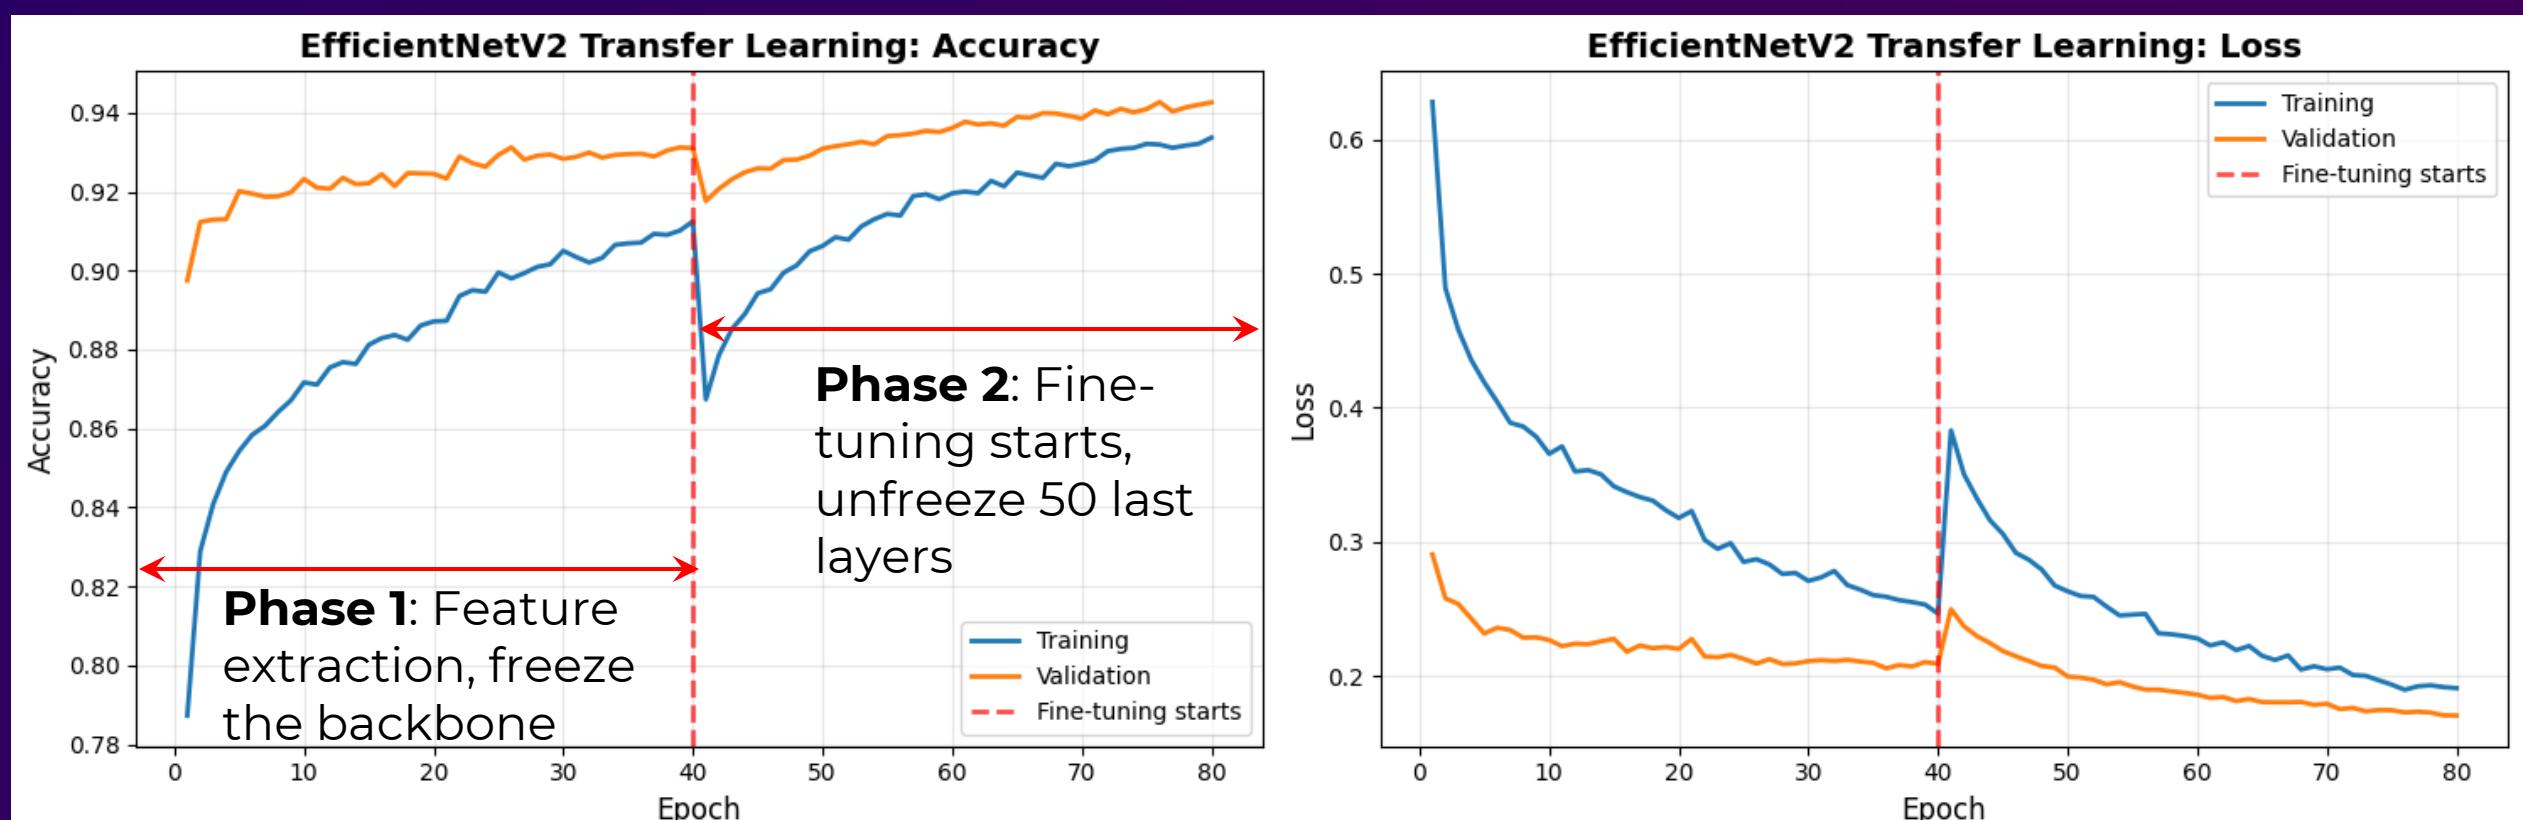
Dense (10) (softmax)

**Total params** 6.26M |  
**Trainable** 334K (Phase 1)

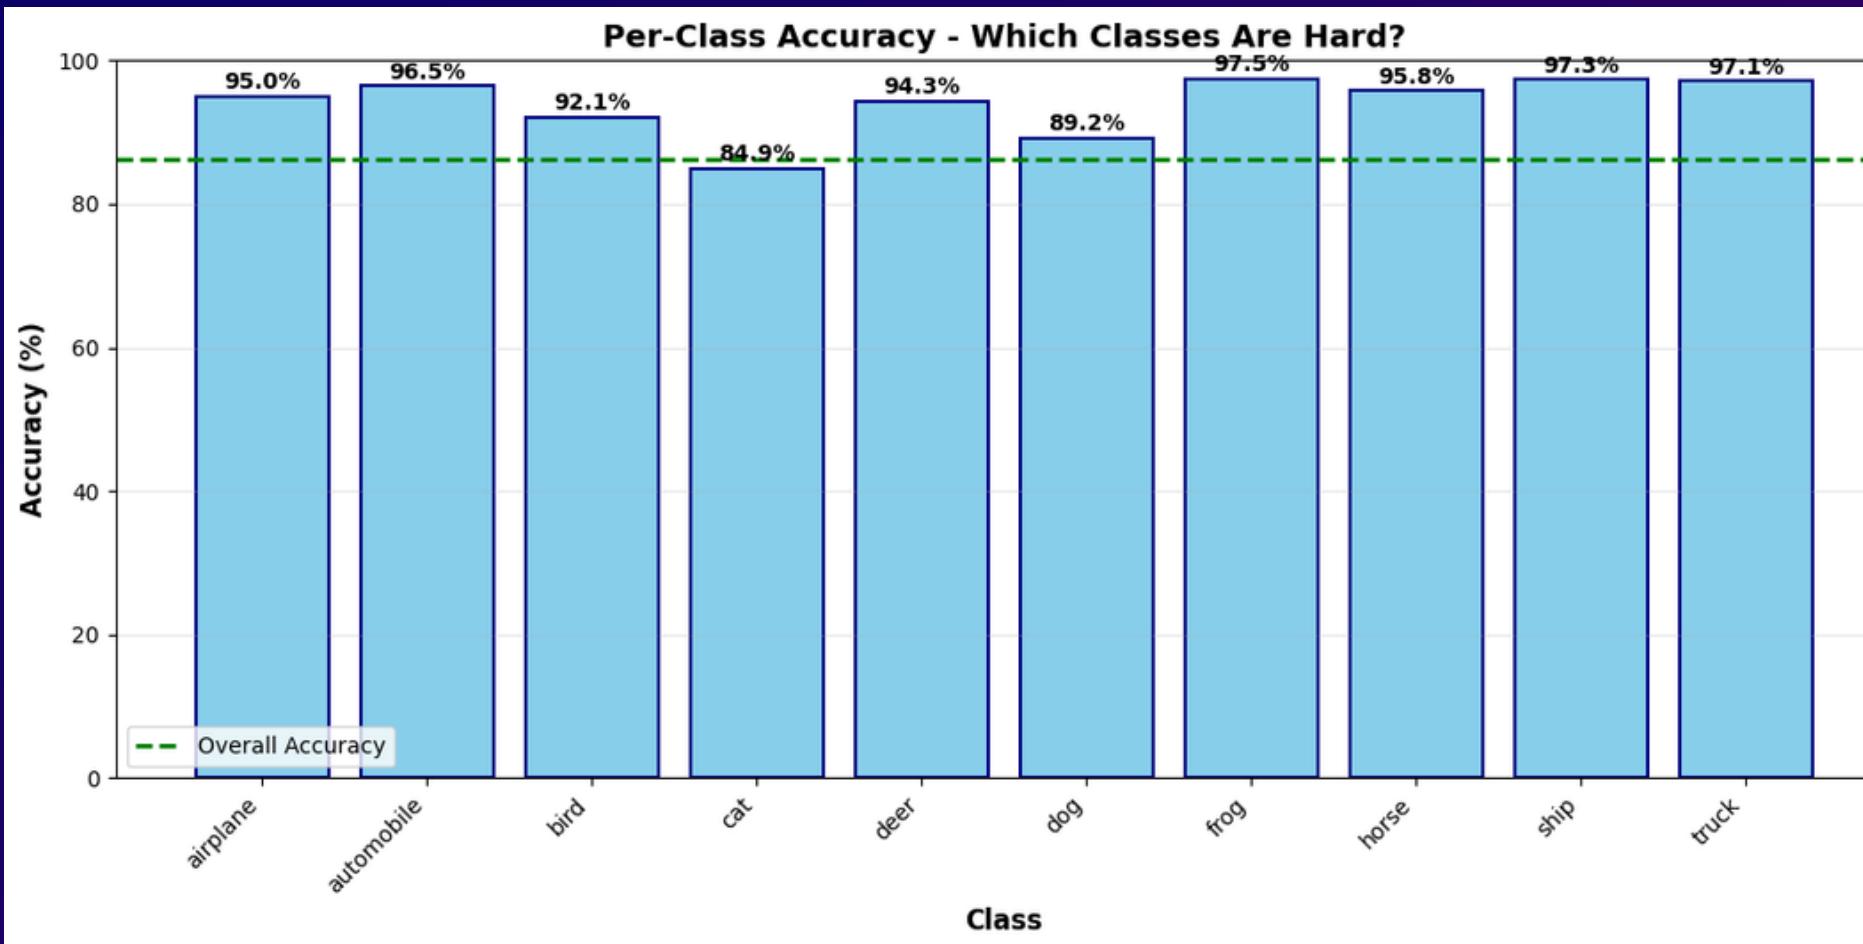
Transfer learning gave a strong jump in accuracy without training from scratch.

# EfficientNetV2B0 CNN Model

Model	Total Parameters	Key Modifications	Test Acc	Test Loss	Overfitting Gap	Epochs
5-layer CNN	316K	Dropout 0.2	67.5%	1.1	+17.8%	10
VGG10-optimized	403K	+ Early Stop, Reduced LR, Data Augmentation	82.1%	0.5	-1.8%	50/50
ResNet-20	275K	All previous, no Dropout	87.7	0.4	+3.1%	114/200
EfficientNetV2B0	6.26M	All+, Input: 96x96, 30 layers unfrozen	90.4%	0.3	-4.5%	30+30
EfficientNetV2B0	6.26M	All+, Input: 224x224, 50 layers unfrozen (25%)	94.0	0.2	0.0%	40+40

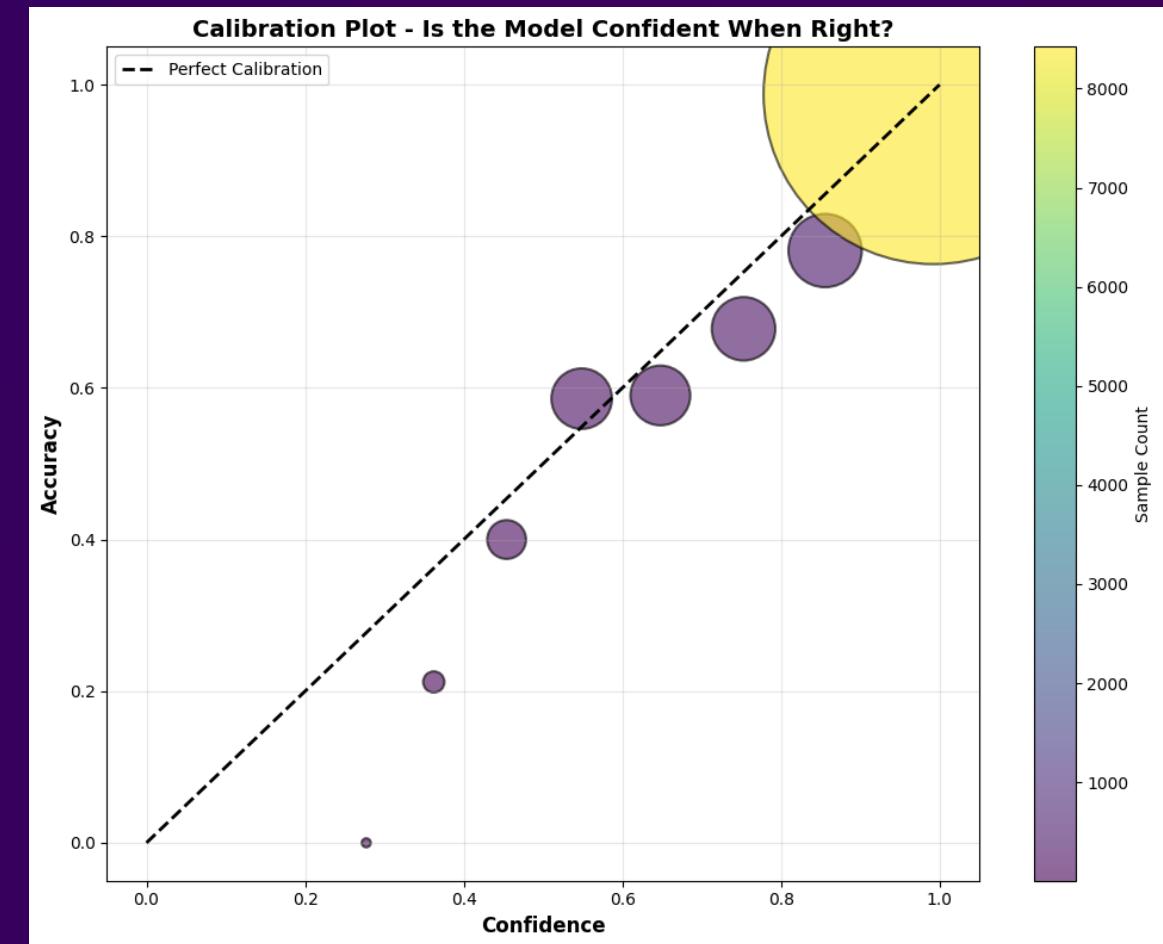


# EfficientNetV2: Evaluation Dashboard

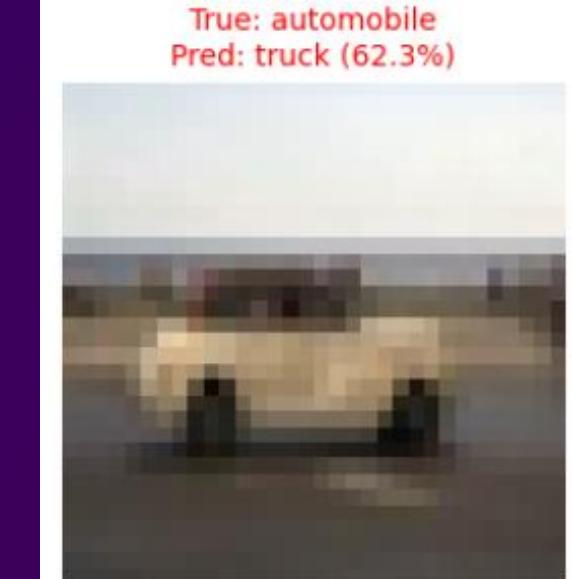
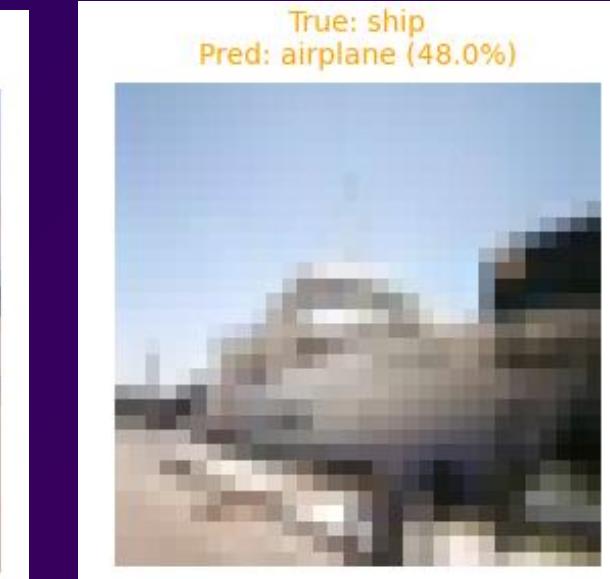
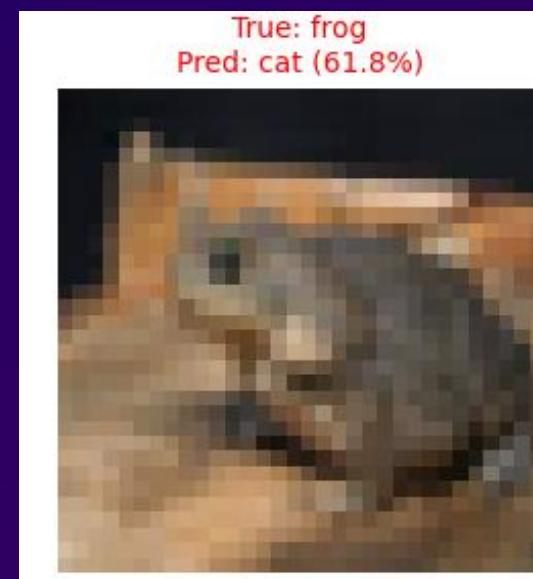
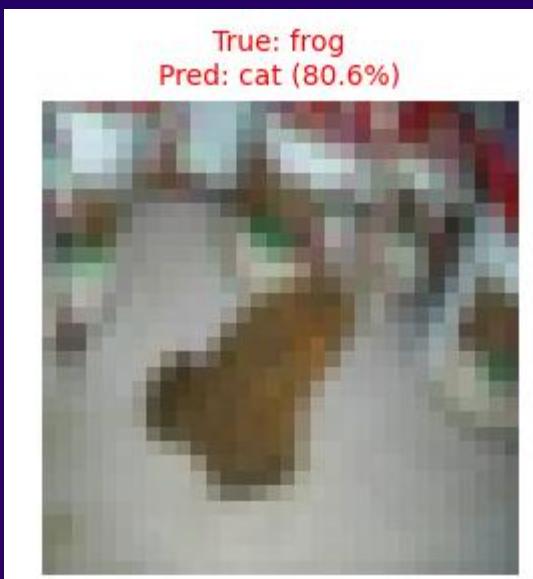


### Top 10 Most Confused Class Pairs:

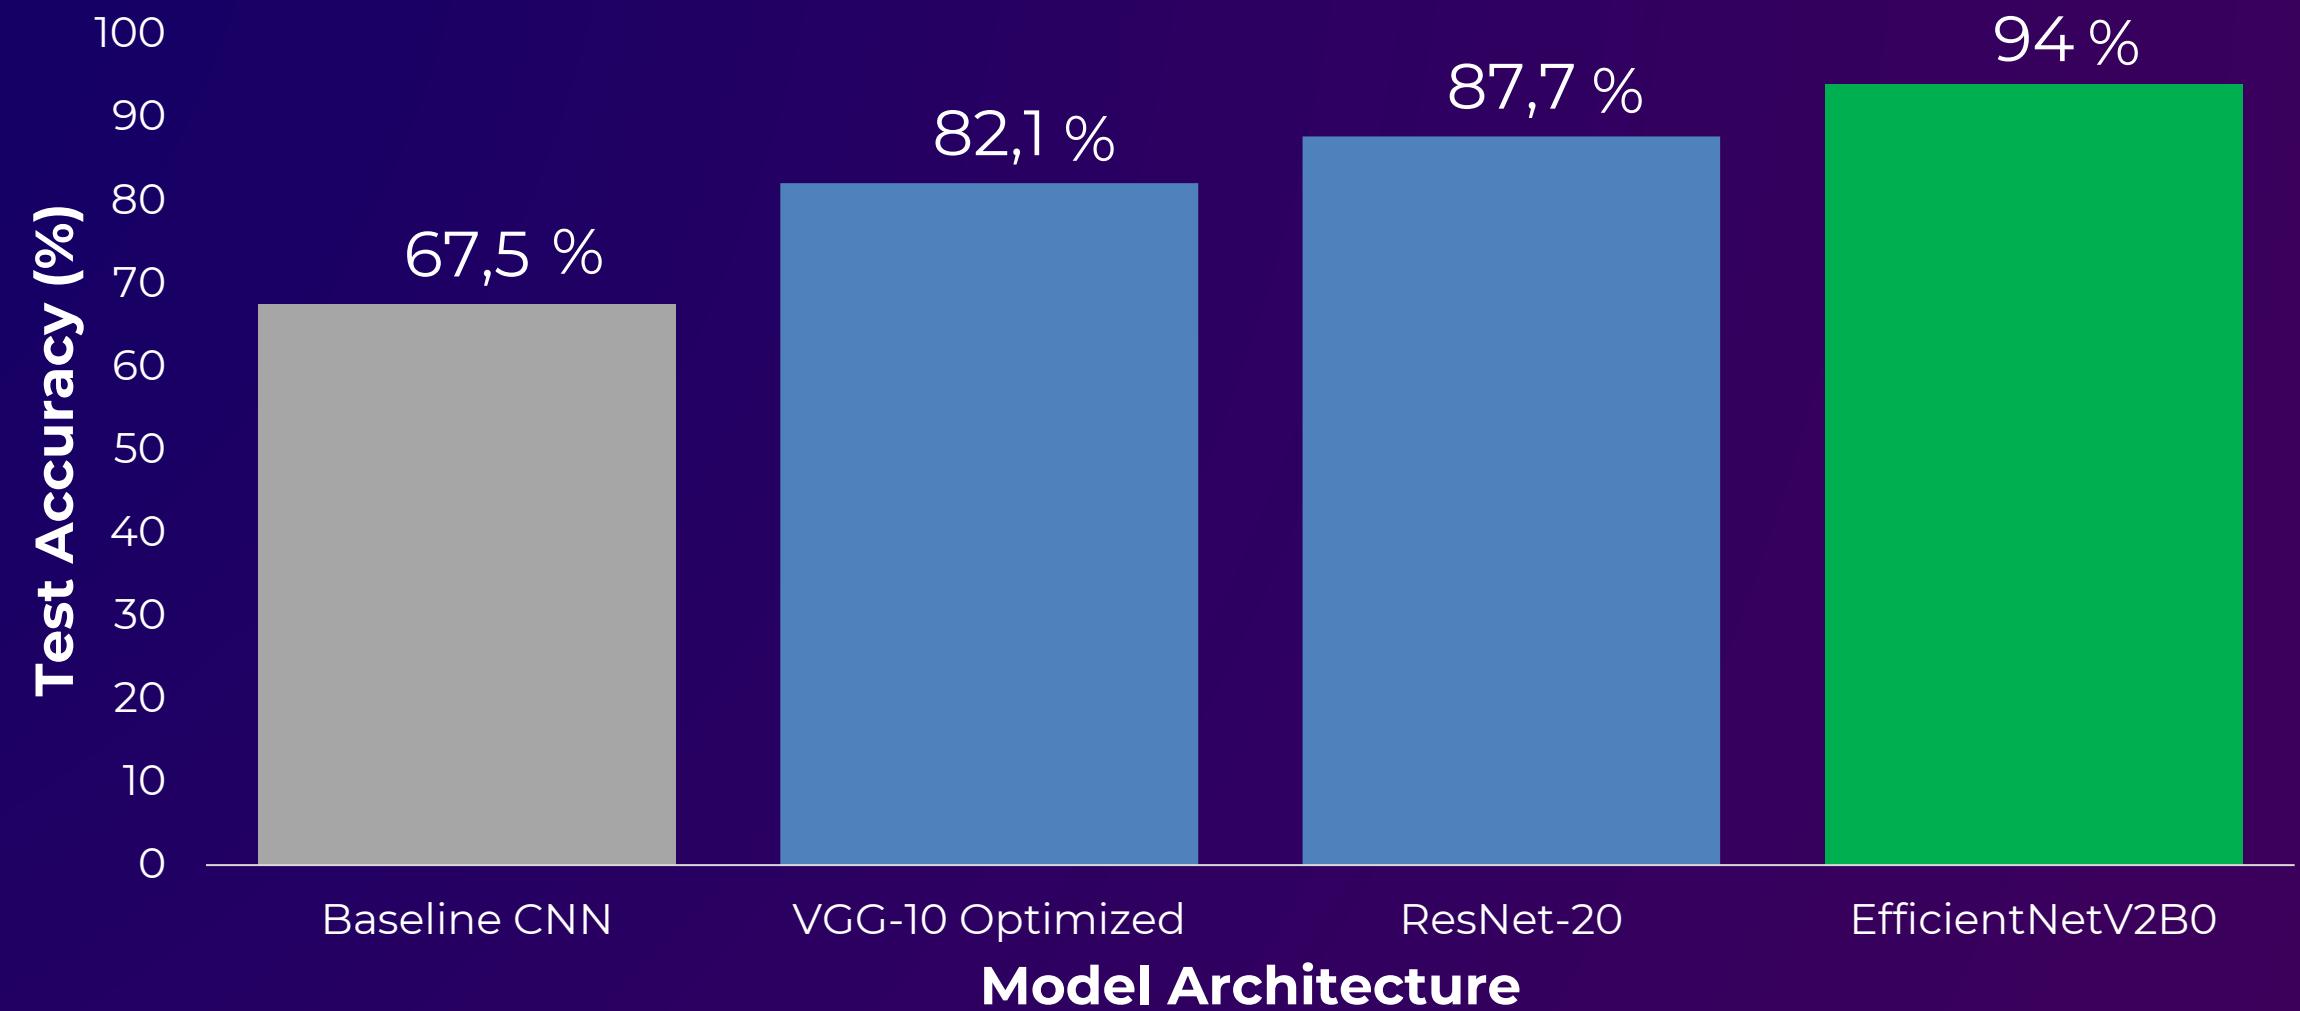
Confusion Pair	Count
Dog→Cat	73
Cat→Dog	70
Automobile→Truck	27
Cat→Frog	25
Cat→Deer	21
Airplane→Ship	20
Bird→Deer	19
Truck→Automobile	19
Deer→Horse	18
Horse→Deer	18



## Misclassified Examples: Where the model went wrong?



# Conclusions



## Best Model

**Best model:** EfficientNetV2B0  
(transfer learning)

**Test accuracy:** 94.0% | **Lowest test loss**

Two-phase training: **freeze → fine-tune (last 50 layers)**

## What improved performance?

**VGG:** deeper conv blocks → big jump vs baseline

**ResNet-20:** skip connections → more stable training

**EfficientNetV2B0:** transfer learning → largest gain with fewer trainable weights

## Error patterns & next steps

Confusion mostly between **similar classes** (dog↔cat, automobile↔truck)

### Next steps :

- stronger augmentation (random crop/ flip etc.,)
- tune fine-tuning depth (unfreeze fewer/more layers)